




The Federated GraphQL Subscriptions Zoo



Tom Houlé

Federated GraphQL



Subscriptions are special... in GraphQL

“a long-lived request that fetches data in response to a sequence of events over time”

— GraphQL spec (draft)

Subscriptions are special... in GraphQL

"a long-lived request that fetches data in response to a sequence of events over time"

— GraphQL spec (draft)

"GraphQL supports type name introspection within any selection set in an operation, with the single exception of selections at the root of a subscription operation."

— GraphQL spec (draft)

Subscriptions are special... in GraphQL

Subscriptions are special... in GraphQL

"Subscription operations must have exactly one root field.

To enable us to determine this without access to runtime variables, we must forbid the @skip and @include directives in the root selection set."

— GraphQL spec (draft)

Subscriptions are special... in GraphQL

"Subscription operations must have exactly one root field.

To enable us to determine this without access to runtime variables, we must forbid the @skip and @include directives in the root selection set."

— GraphQL spec (draft)

"While each subscription must have exactly one root field, a document may contain any number of operations, each of which may contain different root fields. When executed, a document containing multiple subscription operations must provide the operation name as described in GetOperation()."

— GraphQL spec (draft)

Subscriptions are special... in GraphQL-over-HTTP

Subscriptions are special... in GraphQL-over-HTTP

"GraphQL Subscriptions are beyond the scope of this specification at this time."

— GraphQL over HTTP spec (draft)

Subscriptions are special... in GraphQL-over-HTTP

"GraphQL Subscriptions are beyond the scope of this specification at this time."
— GraphQL over HTTP spec (draft)



Subscriptions are actually not that special in Federated GraphQL

Subscriptions are actually not that special in Federated GraphQL

Schema of the sales subgraph:

```
1 type Product @key(fields: "id") {  
2   id: ID!  
3 }  
4  
5 type Subscription {  
6   productSales: Product  
7 }
```



Schema of the products subgraph:

```
1 type Product @key(fields: "id") {  
2   id: ID!  
3   name: String!  
4 }  
5  
6 type Query {  
7   productById(  
8     id: ID!  
9   ): Product @lookup  
10 }
```



Subscriptions are actually not that special in Federated GraphQL

Client → Gateway

```
1 subscription ProductSalesWithName {  
2   productSales {  
3     name  
4   }  
5 }
```



Gateway → sales subgraph

```
1 subscription {  
2   productSales {  
3     id  
4   }  
5 }
```



Gateway → products subgraph

```
1 query {  
2   productById(id: $id) {  
3     name  
4   }  
5 }
```



Subscriptions are actually not that special in Federated GraphQL

Data returned to the client:

```
1  {"name": "Labubu"}
2  {"name": "Labubu"}
3  {"name": "Crocs"}
4  {"name": "Zune"}
5  {"name": "Furbies (12 pack)"}
6  {"name": "Labubu"}
7  {"name": "Google Glass"}
```



The problems with Federated Subscriptions

- Lack of transport standardisation has led to fragmentation:

The problems with Federated Subscriptions

- Lack of transport standardisation has led to fragmentation:
 - WebSockets (HTTP/1.1)
 - Subprotocols with protocol negotiation

```
1 Sec-WebSocket-Version: 13
```

```
2 Sec-WebSocket-Protocol: graphql-ws, graphql-transport-ws
```


The problems with Federated Subscriptions

- Lack of transport standardisation has led to fragmentation:

- WebSockets (HTTP/1.1)

- Subprotocols with protocol negotiation

```
1 Sec-WebSocket-Version: 13
```

```
2 Sec-WebSocket-Protocol: graphql-ws, graphql-transport-ws
```

- Init payloads are not headers

The problems with Federated Subscriptions

- Lack of transport standardisation has led to fragmentation:

- WebSockets (HTTP/1.1)
 - Subprotocols with protocol negotiation

```
1 Sec-WebSocket-Version: 13
```

```
2 Sec-WebSocket-Protocol: graphql-ws, graphql-transport-ws
```

- Init payloads are not headers
- SSE (HTTP/2 and 3)
- Multipart

The problems with Federated Subscriptions

- Lack of transport standardisation has led to fragmentation:

- WebSockets (HTTP/1.1)
 - Subprotocols with protocol negotiation

```
1 Sec-WebSocket-Version: 13
```

```
2 Sec-WebSocket-Protocol: graphql-ws, graphql-transport-ws
```

- Init payloads are not headers
 - SSE (HTTP/2 and 3)
 - Multipart
- One connection between the Gateway and the relevant subgraph per subscribed client, even when they all subscribe to the same events

The problems with Federated Subscriptions

- Lack of transport standardisation has led to fragmentation:

- WebSockets (HTTP/1.1)
 - Subprotocols with protocol negotiation

```
1 Sec-WebSocket-Version: 13
```





```
2 Sec-WebSocket-Protocol: graphql-ws, graphql-transport-ws
```

- Init payloads are not headers
 - SSE (HTTP/2 and 3)
 - Multipart
- One connection between the Gateway and the relevant subgraph per subscribed client, even when they all subscribe to the same events
- Multi-protocol subscriptions

Multi-protocol subscriptions

-  Client —  → Gateway —  →  Subgraph




Multi-protocol subscriptions

-  Client —  → Gateway —  →  Subgraph
- At each step, one of
 - SSE,
 - WebSockets
 - subscriptions-transport-ws
 - graphql-ws / graphql-transport-ws
- And different handshake shapes between each!
 - Headers vs websocket init payload shape mismatch

```
1 interface ConnectionInitMessage {  
2   type: 'connection_init';  
3   payload?: Record<string, unknown> | null;  
4 }
```

TS

Multi-protocol subscriptions

-  Client —  → Gateway —  →  Subgraph
- At each step, one of
 - SSE,
 - WebSockets
 - subscriptions-t
 - graphql-ws / gra
- And different hands
 - Headers vs webs

```
1 interface Connect
2   type: 'connect'
3   payload?: Rec
4 }
```



TS

Alternative: connect the gateway to a message queue

- The idea: the GraphQL federation gateway connects to a message queue (Kafka, NATS, ...), not the subgraphs directly
 - The subgraphs or other services post messages to that queue
- Two implementations
 - EDFS
 - Grafbase extensions

EDFS

```
1  input edfs__NatsStreamConfiguration {
2      consumerInactiveThreshold: Int! = 30
3      consumerName: String!
4      streamName: String!
5  }
6
7  type PublishEventResult {
8      success: Boolean!
9  }
10
11 type Query {
12     employeeFromEvent(id: ID!): Employee!
13     @edfs__natsRequest(subject: "getEmployee.
14         {{ args.id }}")
15 }
16
17 input UpdateEmployeeInput {
18     name: String
19     email: String
20 }
```



```
19
20 type Mutation {
21     updateEmployee(id: ID!, update:
22         UpdateEmployeeInput!): PublishEventResult!
23     @edfs__natsPublish(subject: "updateEmployee.
24         {{ args.id }}")
25 }
26
27 type Subscription {
28     employeeUpdated(employeeID: ID!): Employee!
29     @edfs__natsSubscribe(subjects: ["employeeUpdated.
30         {{ args.employeeID }}"])
31 }
```

Grafbase Extensions

- Pluggable gateway extensions compiled to WebAssembly (WASI preview 2)
 - They can define their own directives that will be used by the Gateway for query planning
 - Near-native performance, in-process secure sandbox.
 - They can perform arbitrary IO (but you can restrict that with permissions).
 - You use the open source extensions on the Grafbase Marketplace or build your own
 - They can act as **virtual subgraphs**



```
1  extend schema
2    @link(
3      url: "https://specs.grafbase.com/composite-schemas/v1"
4      import: ["@lookup", "@key", "@derive"]
5    )
6    @link(
7      url: "https://extensions.grafbase.com/extensions/nats/
8          0.4.1"
9      import: ["@natsPublish", "@natsSubscription"]
10   )
11  input SellProductInput {
12    productId: ID!
13    price: Int!
14  }
15
16  type Mutation {
17    sellProduct(input: SellProductInput!): Boolean!
18    @natsPublish(
19      subject: "productSales",
20      body: { selection: "*" })
21  }
```

```
22
23  type Product @key(fields: "id") {
24    id: ID!
25  }
26
27  type ProductSale {
28    productId: ID!
29    product: Product! @derive
30    price: Int!
31  }
32
33  type Subscription {
34    sales(subject: String!): ProductSale
35    @natsSubscription(
36      subject: "{{ args.subject }}"
37      selection: "select(.price > 10)"
38    )
39  }
```

Corresponding configuration

```
1 [extensions.nats]
2 version = "0.4.1"
3
4 [[extensions.nats.config.endpoint]]
5 servers = ["nats://localhost:4222"]
```

Advantages of an extensions-based approach compared to EDFS

- Arbitrary data formats for the messages (not only JSON)
- Customizable and extensible without touching the Gateway. You can write extensions for other pub/sub systems (Kinesis, etc.).
- More powerful filters (jq expression language)
- By convention, configuration is usually in your Gateway configuration, not expressed in your subgraph's GraphQL schemas

Takeaways

Takeaways

- Federated GraphQL subscriptions require some thinking and planning.

Takeaways

- Federated GraphQL subscriptions require some thinking and planning.
- Pros of traditional federated subscriptions
 - Federate existing GraphQL subgraphs, no need to modify them
 - Subscription fields are managed directly in your subgraphs, next to your other logic

Takeaways

- Federated GraphQL subscriptions require some thinking and planning.
- Pros of traditional federated subscriptions
 - Federate existing GraphQL subgraphs, no need to modify them
 - Subscription fields are managed directly in your subgraphs, next to your other logic
- Pros of subscriptions offloaded to a message queue
 - Stream deduplication
 - Non-GraphQL services can publish to subjects directly
 - Depends on setup, but usually higher performance with less memory usage

Takeaways

- Federated GraphQL subscriptions require some thinking and planning.
- Pros of traditional federated subscriptions
 - Federate existing GraphQL subgraphs, no need to modify them
 - Subscription fields are managed directly in your subgraphs, next to your other logic
- Pros of subscriptions offloaded to a message queue
 - Stream deduplication
 - Non-GraphQL services can publish to subjects directly
 - Depends on setup, but usually higher performance with less memory usage

You can mix and match both approaches

Also

Also

Workshop!

Also

Workshop! Tomorrow!

Also

Workshop! Tomorrow!

Grote Zaal - 2nd Floor.

Also

Workshop! Tomorrow!

Grote Zaal - 2nd Floor. 10:45am.

Also

Workshop! Tomorrow!

Grote Zaal - 2nd Floor. 10:45am.

Thank you!



Appendices

Links

- WebSockets
 - subscriptions-transport-ws
 - Issues and security implications with subscriptions-transport-ws
- SSE
 - GraphQL-SSE spec
- Multipart subscriptions
 - Incremental delivery over HTTP
 - Apollo docs
- Grafbase extensions
- Cosmo EDFS
- Pen Pineapple Apple Pen